Data Quality over Scale: Compact Models for Arabic Diacritization and the Frontier Regression

Interscript ML Team

September 2026

Abstract

Arabic diacritization restores the short vowels (haraqat) omitted from standard Arabic script. The strongest published dedicated system, Sadeed, scores 7.29% DER (with case endings) on its own SadeedDiac-25 benchmark using a 1.5B-parameter decoder-only language model. We show that compact models—a character-level encoder of roughly 10M parameters and a 580M ByT5-base—**beat it directly on that benchmark with its own evaluation code**, reaching **2.29% DER (CE) and 1.33% without case endings** under a zero-skip windowed protocol: 3.2× and 4.0× better than Sadeed, ahead of GPT-4 (3.86%) and Gemini-Flash-2.0 (3.19%/2.38%). We additionally measure the LLM frontier on this benchmark and find a family-wide regression: GLM-5.2 scores 2.51%/2.69% (raw/zero-skip) under a neutral protocol—vowel knowledge matching our dedicated 580M teacher to within 0.02pp—while every GLM successor we measured is 3–5× worse (5.3-Flash 8.80, 5.3 9.90, 4.7-flash 13.23 zero-skip), with per-position attribution locating the loss in wrong haraqat rather than writing convention. A distilled 300M client-tier student reaches 4.82% (42% error reduction) through a controlled lever ladder whose two registered negatives—multi-token-prediction auxiliary and constant-budget register mixing—mirror the RL negative: at SFT convergence every verified gain came from supervision quality, not from policy optimization, frontier training techniques, or register diversification. Our combined corpus merges the full 75M-word Tashkeela dump (cleaned with a Sadeed-style pipeline), Arabic Wikipedia, and QCRI EMNLP-2025 data, totalling 2.1M sentences; scaling 28× (75K → 2.1M) reduced in-domain DER from 2.42% to 0.99%. We further show that Misraj's public training corpus *contains its own benchmark* (122 paragraphs verbatim plus ∼1k near-duplicates) and decontaminate it before domain adaptation. We argue that for well-resourced morphological restoration tasks, data curation dominates model capacity, and we release all training and evaluation code.

1 Introduction

Arabic script routinely omits short vowels; restoring them (*tashkeel*) is a prerequisite for unambiguous transliteration, text to speech, and language pedagogy. Prior work has escalated model size: Sadeed [1] fine-tunes a 1.5B-parameter Arabic language model (1.2% DER on its internal test set; 7.29% on the harder SadeedDiac-25 benchmark it introduced).

We take the opposite direction. Our contributions:

1. A compact char-level encoder (6 layers, $d = 384$, 6 heads, ≈10M parameters) that **beats Sadeed on SadeedDiac-25 with Misraj's own evaluator**: 3.25% DER (CE) / 1.81% (w/o CE) vs 7.29% / 5.26%, at 1/150th the parameters; a 580M ByT5-base [2] extends this to 2.81% / 1.69%, beating Gemini-Flash-2.0 on all four metrics under a zero-skip windowed protocol.

2. A data ablation showing in-domain DER falls 2.42% → 0.99% as the corpus grows 75K → 2.1M (28×).

3. A contamination audit of the benchmark's public source corpus (122 verbatim + ∼1k near-duplicate lines), a stride-1 60-character-window decontamination, and a zero-skip windowed evaluation protocol.

4. Paragraph-context training (joining corpus lines into ∼1400-byte units) closing the context gap to LLMs: 2.81→2.68% DER (CE), 1.69→1.60% (w/o CE).

5. A teacher lineage driven by supervision quality, not architecture: a morphological-analysis auxiliary stream (r6, 2.5793% DER) and a teacher-labeled news-domain mix (r7, **2.2864% / 1.3343%**)—the best dedicated model measured on this benchmark—with a controlled ablation attributing the r6 gain to lexical-morphological knowledge injection (IPA-aux 2.6588 vs none 2.6775 vs morph 2.5793).

6. A frontier regression axis: under our neutral protocol, GLM-5.2 scores 2.69% zero-skip DER—and every GLM successor we measured is 3–5× worse (5.3-Flash 8.80, 5.3 9.90, 4.7-flash 13.23), with per-position attribution locating the loss in wrong haraqat (10.05% of positions vs 5.2's 2.64%, which matches our 580M teacher's 2.62%), not writing convention (the dagger-alif effect is 0.125pp).

7. A distilled client tier at 300M parameters: 8.26% → **4.82%** (42% error reduction at identical architecture) via a controlled lever ladder (optimizer −2.96pp, fresher teacher labels −0.47pp, longer training −0.25pp), plus four registered negatives (multi-token-prediction auxiliary +0.26pp; register swap +0.98pp; register add flat-with-epochs-cancelled at 4.82; on-policy GKD +1.18pp, the worst rung) showing that neither frontier-training techniques, register mixing in either direction, nor on-policy exposure substitute for supervision quality.

8. A three-way RL negative result (RAFT, sequence-GRPO on Persian, entropy-weighted GTPO-GRPO on Arabic): all flat or negative at SFT convergence — diacritization residual error is knowledge-limited, not policy-limited.

2 Related Work

Sadeed [1] is the strongest published Arabic diacritizer (1.5B params; 1.2% DER in-domain, 7.29% on SadeedDiac-25). Earlier systems include Shakkelha, Tafran, and CTC-based encoders; all are dominated by both Sadeed and this work.

3 Data

3.1 Sources

- **Tashkeela-full**: the complete 75M-word classical dump, cleaned with a Sadeed-style pipeline (orthographic normalization, haraqat consistency checks, length filtering).
- **Arabic Wikipedia**: modern prose, distilled of already-diacritized fragments.
- **QCRI EMNLP-2025**: news-domain diacritized data.

Combined: 2.1M sentences (train), 240K (val), 52,224 held-out test.

3.2 Cleaning

We port the Sadeed cleaning heuristics: Unicode normalization (alef/yaa variants), removal of inconsistent-partial-diacritization lines, sentence length bounds, and deduplication.

4 Model

Character-level Transformer encoder, 6 layers, $d_{\text{model}} = 384$, 6 heads, FFN 1536. Input is the undiacritized character sequence; each consonant position emits a multi-class head over haraqat combinations. Training: cross-entropy with label smoothing 0.1, cosine schedule, 15 epochs, batch 64.

5 Experiments

5.1 Data scaling

Corpus	Sentences	DER
Initial (Tashkeela sample)	75K	2.42%
Combined v2 (full)	2.1M	**0.99%**

Table 1: 28× data scaling halves the error rate.

5.2 Comparison on SadeedDiac-25 (their benchmark, their evaluator)

5.3 Paragraph-context training

SFT units were isolated ≤640-byte lines while the benchmark and the LLMs read whole paragraphs. The decontaminated Misraj corpus is line-split continuous book text, so joining k adjacent lines reconstructs true paragraphs. Training r3's best for one epoch on ∼1400-byte joined units (500K domain + 100K replay) closed the context gap: 2.81→2.68% DER (CE), 1.69→1.60% (w/o CE), overtaking the verified frontier (GLM-5.2 zero-skip 2.69%/1.72%) on both metrics. The same specialization costs +0.53 WER on the out-of-domain WikiNews-2024 multi-reference probe.

5.4 The frontier regression axis

The paragraph-context result above overtook GLM-5.2—but the frontier moved underneath us. Measuring the GLM family under the identical protocol (same 1,200 paragraphs, same evaluator, temperature 0, plain completion where the API still allows it): GLM-5.2 reproduces at 2.5060% raw / 2.6911% zero-skip; GLM-5.3-Flash lands at 8.5721 / 8.7978; the full GLM-5.3 at 9.9760 / 9.8971; and glm-4.7-flash, the last plain-completion GLM model, at 13.0035 / 13.2256 (12 resumable passes past sustained provider rate limits; all 1,200 responses real). Per-position attribution over the same paragraphs (convention-normalized, rates over the 179,401 ground-truth-marked positions) localizes the loss: GLM-5.2's wrong-haraqat rate is 2.64%—matching our dedicated r7 teacher's 2.62% to within 0.02pp—while 5.3-Flash's is 10.05% and 4.7-flash regresses on both axes at once (wrong 9.01%, missing 6.67%, the family's worst under-diacritization). The entire dagger-alif (U+0670) convention effect is 0.125pp. Paired-bootstrap CIs put every successor decisively behind 5.2 (+7.9 to +8.0pp, CIs excluding zero by a wide margin). As frontier optimization has shifted agentic, the

System	Params	DER (CE)	DER (w/o CE)	WER (CE)
Claude-3-7-Sonnet	—	1.3941	0.7693	4.6718
Ours (ByT5-base r7, morph-aux + news mix)	580M	**2.2864**	**1.3343**	—
Ours (ByT5-base r6, morph-aux)	580M	2.5793	1.5317	—
GLM-5.2 (our reproduction, raw)	—	2.5060	1.5537	7.9929
GLM-5.2 (our reproduction, zero-skip)	—	2.6911	1.7179	8.3037
GLM-5.3-Flash (zero-skip, reasoning_effort=low)	—	8.7978	6.6368	30.84
GLM-5.3 (zero-skip, reasoning_effort=low)	—	9.8971	7.8219	31.07
glm-4.7-flash (zero-skip, thinking-disabled)	—	13.2256	10.3206	36.10
Gemini-Flash-2.0	—	3.1926	2.3783	7.9942
GPT-4	—	3.8645	3.8645	5.2719
Sadeed	1.5B	7.2915	5.2625	13.7425
Ours (ByT5 r5 paragraph-context, windowed)	580M	2.6775	1.5965	8.0919
Ours (r5 + GTPO-GRPO, 700B windows)	580M	2.6597	1.5818	8.1165
Ours (ByT5 r3, windowed)	580M	2.8126	1.6877	8.4110
Ours (ByT5 r3, single-shot)	580M	2.8429	1.7589	8.4981
Ours (r3 + RAFT, beam 4 / beam 1)	580M	2.8515 / 2.8308	1.7617 / 1.7638	—
Ours (encoder)	≈10M	3.2495	1.8072	10.3276

Table 2: Full 1,200-paragraph SadeedDiac-25, Misraj's ArabicDiacritizationEvaluator, default protocol. Published LLM/Sadeed rows from the benchmark card; the GLM-5.3-generation rows are our reproductions (thinking cannot be disabled on 5.x—HTTP 400 code 1210—so they run at `reasoning_effort=low`, the nearest expressible analog of plain completion; glm-4.7-flash still accepts disabled thinking). The windowed protocol splits long inputs at word boundaries into ≤600-byte in-distribution windows with a 2× generation cap (diacritized output is 1.4–1.6× input bytes; a naive input-length cap silently truncates) and projects haraqat onto input letters, so *zero* paragraphs are skipped—a strictly cleaner protocol than the published rows.

newest generalists lost classical-Arabic mark knowledge their predecessor had—while our dedicated 580M teacher improved.

5.5 The distilled client tier

The same benchmark disciplines a 300M distilled student shipped for browser/edge inference. From the 1.0 rung (8.26%), a pre-registered lever ladder reached **4.82%** (42% error reduction at identical architecture and artifact size): the Muon optimizer [3] alone contributes −2.96pp (a controlled 2×2 factorial closes additively: optimizer −2.96, memory-layer capacity −0.70, combined −3.43), fresher r7 teacher labels −0.47pp, and longer training a further −0.25pp (4.57%). Four registered negatives bracket the ladder: multi-token-prediction [4] as a training auxiliary scored 5.09% (+0.26pp vs control, with a disclosed preemption confound); swapping news-domain training units for classical Tashkeela at constant budget scored 5.81% (+0.98pp); adding 18k classical units on top at matched epochs scored 4.82%—flat against the 3-epoch control and 0.25pp behind news-only training at the same epochs, the add canceling the epoch gain entirely; and on-policy generalized distillation [5] (reverse-KL on student-sampled sequences scored by the frozen teacher) scored 6.00% (+1.18pp, the worst rung measured)—the domain-coverage hypothesis failing in every direction it was tested. Together with the RL negatives, the pattern is one line: at SFT convergence, supervision quality dominates policy optimization, frontier training techniques, on-policy exposure, and register diversification alike. The tier ships as checksummed, runtime-agnostic artifacts (IMF v1) whose quantization fragility was itself diagnosed and fixed—the quantized output head, not the body, caused confident flip errors (9.34% → 0.26% after the head-fp32 repair).

5.6 RL at SFT convergence: three negative results

We tested whether the residual (98.7% phonotactically legal alternates) can be sharpened by policy optimization:

- **RAFT** (rejection-sampling FT, K=8, 3 iterations, 1,170 winners): benchmark flat — 2.8429→2.8515 (beam 4) / 2.8308 (beam 1).
- **GRPO** (sequence reward, G=8, KL leash): on Persian G2P the best checkpoint *lost* 2 points on SentenceBench homographs (77.34→75.37%), with dev reward degrading monotonically past step 200.
- **GTPO-GRPO** (entropy-weighted per-token credit [6]; graded alignment-based reward): dev curve *exactly* flat (5.9692% at steps 0/100/150); the final benchmark equals the teacher within protocol noise.

Conclusion: at SFT convergence on 2M clean lines, the residual is **knowledge-limited, not policy-limited**. Every lever that moved the benchmark was data-side: decontamination (−0.1 DER), domain adaptation (−0.1), paragraph context (−0.14).

6 Discussion

Data curation dominates capacity. A 150× smaller encoder beats a 1.5B model on its own benchmark once the corpus is cleaned and scaled, and the 580M ByT5-base lineage reaches 2.29% DER (CE) / 1.33% (w/o CE) through supervision-side moves alone (paragraph context, morphological auxiliary, teacher-labeled domain mix)—beating Gemini-Flash-2.0 on all four metrics and every GLM generation after 5.2 on every metric. On our in-domain split the models reach ≤1.4% DER; the gap to 2.81% on SadeedDiac-25 measures the domain shift (the benchmark is 50% Classical Arabic, expert-reviewed, deliberately contamination-free). Error analysis shows the residual is 67% word-internal vowel confusion and 33% word-final case endings: domain knowledge, not içrāb structure.

The public corpus leaks its own benchmark. Misraj's released training corpus (1.88M lines) contains 122 benchmark paragraphs verbatim (diacritics-stripped match) plus ∼1k lines sharing 60+-character windows with benchmark paragraphs. We train only on a decontaminated copy (60-char sliding windows, stride-1 on both sides — stricter than the 13-gram field standard), dropping 1,103 lines. Any number trained on the raw release would be contaminated.

Protocol discipline. An intermediate evaluation of ours scored 1.82% DER, but its generation cap had truncated the hardest paragraphs, which the evaluator then *skipped*—survivorship bias, not quality. The zero-skip windowed protocol above is the only number we report.

7 Limitations

Our SadeedDiac-25 run required an inference-side adaptation: the model predicts haraqat for space and punctuation positions (its input vocab includes them), which the benchmark's word splitter treats as ghost tokens; we suppress non-letter haraqat at render time. The published Claude-3.7-Sonnet row (1.39%) is vendor-published under an undisclosed protocol; we do not treat it as reproducible and note only that a newer flagship measured under our fully published protocol scores 2.51% raw. WER (CE) trails GPT-4 and Gemini—LLMs copy words more faithfully at the cost of 2× worse DER.

8 Reproducibility

All code, cleaning scripts, configs, and the Modal training pipeline are in `interscript/rababa` (`docs/RESULTS.md` for exact run IDs).

References

[1] Sadeed: Arabic Diacritization. arXiv:2504.21635, 2025.

[2] Xue et al. ByT5: Towards a token-free future with pre-trained byte-to-byte models. TACL, 2022. arXiv:2105.13626.

[3] Liu et al. Muon is Scalable for LLM Training. arXiv:2502.16982, 2025.

[4] Gloeckle et al. Better & Faster Large Language Models via Multi-token Prediction. arXiv:2404.19737, 2024.

[5] Agarwal et al. On-Policy Distillation of Language Models: Learning from Self-Generated Mistakes. ICLR, 2024. arXiv:2306.13649.

[6] Tan et al. GTPO and GRPO-S: Token and Sequence-Level Reward Shaping with Policy Entropy. arXiv:2508.04349, 2025.